

# The Concepts of Kubernetes

## - *Configuration*

Speaker: hp.huang

# Outline

- Best Practices
- Managing Compute Resources for Containers
- Assigning Pods to Nodes
- Evicting Pods from Nodes
- Managing Sensitive Information
- Organizing Clusters Access
- Pod Priority and Preemption

# Best Practices - General Configuration Tips

- Specify latest stable API version
- Store the configuration with version control (GitOps)
- Write the configuration file using **YAML** rather than JSON
- DON'T specify default value
- Put object descriptions in **annotations**
- Use **kubectl**
  - execute command at **directory** layer
  - Use **label selectors** for **kubectl get**, **Kubectl delete**
  - Use **kubectl run**, **kubectl expose**

# Best Practices - Others

- DON'T use ***Naked Pods***
  - they will not be rescheduled in the event of a node failure.
- Create a Service before the workload is going to access it
- Define and use ***labels*** that identify semantic attrs. of application or Deployment
  - *{ app: myapp, tier: frontend, phase: test, deployment: v3 }*

# Managing Compute Resources for Containers

- Resource Types
  - CPU, Memory
- Resource ***requests*** of Node
- Resource ***limits*** of each container

# Managing Compute Resources for Containers

## Troubleshooting

- Give a scenario
  - The Pods are pending with event message *failedScheduling*

```
kubectl describe pod frontend | grep -A 3 Events
```

```
Events:
  FirstSeen LastSeen  Count  From              Subobject    PathReason    Message
  36s      5s      6      {scheduler }      FailedScheduling  Failed for reason PodExceedsFreeCPU and possibly others
```

- Actions
  1. Add more nodes to the cluster
  2. Terminate unneeded Pods to make room for pending Pods
  3. Check that the Pod is not larger than all the nodes

# Assigning Pods to Nodes – Node Selector

- To constrain/prefer pods to run on particular nodes
- The recommended approach – use ***label*** selectors
  - Attach ***label*** to the node
  - Add a ***nodeSelector*** field to the pod configuration

# Assigning Pods to Nodes – Use Case

- The service needs to run on specific IPs



# Evicting Pods from Nodes - Taints and Tolerations

- To ensure that pods are not scheduled onto inappropriate nodes

- **Taint**

- *kubectl taint nodes <node\_name> <key>=<value>:<Effect>*
- for example:

```
kubectl taint nodes node1 key=value:NoSchedule
```

- **Toleration**

```
tolerations:  
- key: "key"  
  operator: "Equal"  
  value: "value"  
  effect: "NoSchedule"
```

```
tolerations:  
- key: "key"  
  operator: "Exists"  
  effect: "NoSchedule"
```

- **Effect**

- *NoSchedule*
- *PreferNoSchedule*
- *NoExecute*\*

# Evicting Pods from Nodes – Use Case

- Dedicated Nodes\*
  - *kubectl taint nodes nodename **dedicated=groupName:NoSchedule***
- Nodes with Special Hardware
  - *kubectl taint nodes nodename **special=true:NoSchedule***
  - *kubectl taint nodes nodename **special=true:PreferNoSchedule***

\*The pods with the tolerations will then be allowed to use the tainted (dedicated) nodes as well as any other nodes in the cluster

# Managing Sensitive Information - Secrets

- Creating Secrets
  - Creating Secrets Using *kubectl*
  - Creating Secrets Manually
  - Creating Secrets Using *Generator*
- Using Secrets
  - Using Secrets as *Files* from a Pod
  - Using Secrets as *Environment Variables*

# Organizing Clusters Access - kubeconfig

- Organize information about ***clusters, users, namespaces***, and ***authentication mechanisms***.
- The consist of Context
  - cluster
  - namespace
  - user
- Partition & Merging
  - <https://github.com/hp-huang-tw/kubernetes/tree/master/config-exercise>

# Pod Priority and Preemption

- **Priority**

- The Pod can has priority. [Stable from kubernetes 1.14]

- ***PriorityClass***

```
apiVersion: scheduling.k8s.io/v1
kind: PriorityClass
metadata:
  name: high-priority
value: 1000000
globalDefault: false
description: "This priority class should be used for XYZ service pods only."
```

- ***PriorityClassName***

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    env: test
spec:
  containers:
    - name: nginx
      image: nginx
      imagePullPolicy: IfNotPresent
      priorityClassName: high-priority
```

- **Preemption**

- Give an example
    - There's a pending Pod P which priority is 1000, there're Pods  $Q_{1-n}$  which priority is 1 and is running on Nodes  $N_{1-m}$ .

# Reference

- <https://kubernetes.io/docs/concepts/configuration/>